

AGGREGATOR MANUAL

Author: Andrej Gajduk

Contents

Data flow of a typical aggregator.....	1
Getting started.....	2
DataReader.....	2
Feature Extractor.....	5
Writing your own feature extractor.....	6
Classifiers.....	7
Clustering.....	8
CategoryMap.....	4
Querying	9

1. Data flow of a typical aggregator

Three different components which are executed at different times and with different frequency are the core of every aggregator. First, the initial run fetches most of the data out on the web and does some initial processing. Second, the incremental run which is executed e.g. every 24 hours gathers any newly published data and processes it. The third component is related to presenting the data to the user in your web application. These three components are made up of different tasks as detailed in the Table below, where you will also find the relevant section for each task.

No.	Main task	Relevant section
1. Initial batch processing		
1	Crawl data from web	DataReader
2	Define the relevant categories and build the Category Map	CategoryMap
3	Extract features on the gathered data	FeatureExtractor
4	Train a classifier for each set of categories	Classifiers
2. Incremental processing		
1	Crawl new data	DataReader
2	Extract features on the new data	FeatureExtractor
3	Classify any new features that have missing information	Classifiers
4	identify duplicate entries by clustering	Clustering
3. Front-end		
Get relevant search options/categories		CategoryMap
Get data that fits a specified search criteria		Querying
How to represent near-duplicate entries in your search results		Clustering

2. Getting started

First you need to decide the domain of interest, e.g. scholarship applications listings or news. You should find some relevant websites that are easy to access, no login required, no AJAX just pure HTML. Next you should write an implementation of `DataReader` that extracts relevant information from the html source. Data readers are web-site specific, meaning that you will probably need to write a new data reader for every site you wish to add to your aggregator. Best to start with a few sites, then add new ones with time. Next fetch any data you can find with your data readers then store in a database using the `DatabaseManager` class.

In the database you now have data amassed from different websites, where data is organized differently. For example one website may tag their scholarships with “fields of study” such as economics, law, medicine, technology etc. while another may have them categorized by their “purpose” like study or research. In the aggregator this is defined in a `CategoryMap` where you can specify all the possible categories and sub-categories that users will be able to browse and search on. For example, a news site will have sports and politics in tier 1 with sports further diversified to football and tennis in tier 2. Naturally not all data entries have the same labels on every website. This is where pattern recognition comes in. You will use data analysis tools (Weka) instead of manually tagging each data object in your database.

To use Weka you first need to perform feature extraction on the data. You can find out more about feature extraction on the web. To extract useful features in the aggregator you will need a `FeatureExtractor` which will process the data you have stored in the database and produce an `Instances` object that you can use directly with Weka. The aggregator comes with an implementation of the term-document-frequency metrics, ready to be applied to any text data.

When you have an `Instances` object you can use it to train a classifier that you will use later to distinguish between sport news and political news. Note that when training your classifier you can only use data that has the topic (sport/politics) information fetched from the website. Some websites may not have this information, and they need to be filtered out. After you have trained your classifier you can use it to classify any data that is missing its topic (sport/politics).

Another important thing you have to consider is the grouping of near-duplicate data entries. This is done by clusterers. For example you might use a clusterer to bundle together all listings of one same scholarship that is published on different websites.

3. DataReader

This is the part where you have to do all the work because every site is its own story. Once you decided upon a site e.g. `www.MySite.com` you need to create a new class that extends `DataReader` and implements the abstract methods

```
void init(int max_links) : find all the relevant links and put them in links  
Data getDataForLink(String link) : extract any relevant information on a link
```

For the `init` function you will need to find a URL that lists all links to the data featured on that website. Usually an empty search or a browse button will do the trick. If there is a paging system you can either set the page size to a large number and get all the links at once or loop through all

the pages and fetch them in groups. Note that all data readers are required to make use of the parameter *max_links*. This just means that as you're populating the links list you should limit yourself to this number, any extra links won't be processed so you might as well save your time.

```
public void init(int max_links) {
    html =
    UrlFetcher.fetchGet("http://www.mysite.com/scholarships/search/?start=0&length=10000");
    //we get all the link in one GET
    Matcher matches = getLinksPattern().matcher(html); /*the getLinksPattern function is
    specific to every site*/
    while ( matches.find() && links.size() < max_links ) { /*we respect the max_links
    parameter*/
        String link = matches.group(1); /*we get a valid url i.e.
        "http://www.mysite.com/scholarships/browse?id=124156" */
        links.add(link); /*we add the URL to the list
    }
}
```

To implement the `getDataForLink` function you should identify what information is relevant for you and what would like to extract. A good practice is to extract everything that seems remotely interesting even if currently you have no need for it. I recommend using the [Pattern](#) class provided from Java. Always try to extract small elements from the page and organize them in a logical way. Don't just have one huge content element. To clarify what this means let us suppose the site provides the following information, taken from the scholarship example:

Summary:

Duration: 1 month

Purpose: Study

Level: Bachelor

Benefits: travel expenses covered, student dorm accommodations

Description:

The University of "some university" offers these fine scholarships to all....

lengthy text ... looking forward to next semester.

Instead of just having the contents split in summary and description you should extract the duration, purpose, level and benefits as separate elements. If you want you can nest them with summary so you keep the initial structure. In this step you can also specify what information will be used by the classifiers, although this can also be done later. The code for the data reader would look something like:

```
public Data getDataForLink(String link) {
    String duration = null;
    Matcher duration_matcher = duration_pattern.matcher(html);
    if ( duration_matcher.find() ) {
        duration = duration_matcher.group(1);
    }
    //repeat for every pattern
    //...
    //create a new TextData object
    TextData text_data = new TextData(title, source, new Date(), application_deadline);
    //create the content, layer by layer
    Content summary = new Content("summary");
    summary.addContent("duration", duration);
    //...
    //add content to the data object
    text_data.addContent(summary);
    //set the main content (the largest piece of text)
```

```

        text_data.setMainContent(description); /*it's very important to set something as the main
content, because by default some of the feature extraction methods use only the main content*/
        //set which fields can be classified upon
        text_data.addCategoryInformaion("level",level);
        return text_data;
    }

```

The `DataReader` is organized as a Java iterator meaning it should support the Boolean `hasNext()` and `Data next()` functions. These functions are already implemented to work with the [links](#) variable list so you don't need to override them. Now that you have written your data readers using one is straightforward.

```

DatabaseManager databasemgt = DatabaseManager.getInstance();
DatabaseManager DataReader reader = new MySiteTextDataReader();
reader.init(1000);
while(reader.hasNext() ) {
    Data data = reader.next();
    databasemgt.addDataToProblem(data,"scholarships");
}

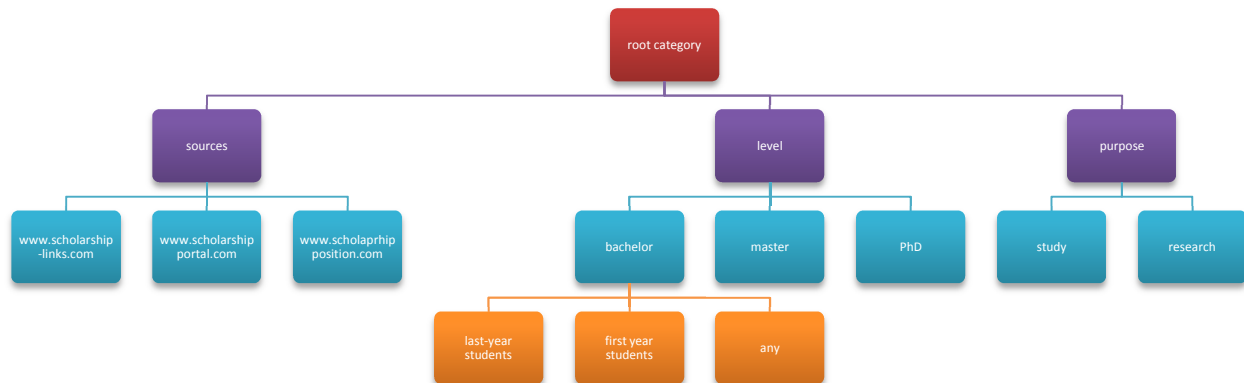
```

If you wanted to check if the data already exists in the database you would use the method:

`databasemgt.checkDataForProblem(data,"scholarships",compare_sources)` - the last argument should be true if you are only interested in matching with data from the same source (website).

4. CategoryMap

The category map is a very important part of our framework. It defines the logical way data is separated. Let's see an example of a category map for our scholarships project.



You can create this `CategoryMap` with the `CategoryMapBuilder`

```

CategoryMap cm = new CategoryMap();
CategoryMapBuilder builder = new CategoryMapBuilder();
builder.addClassToCategoryMap(cm,data,"level");
builder.addClassToCategoryMap(cm,data,"bachelor","level");
builder.addClassToCategoryMap(cm,data,"purpose");
builder.addSourcesToCategoryMap(cm,data);

```

You would normally do this after your initial batch of data crawling. Saving and loading from database is done by:

```

databasemgt.saveCategoryMapForProblem(cm,"scholarships");
databasemgt.getCategoryMapForProblem("scholarships");

```

The CategoryMap is useful when querying on the database, as detailed in Section 9 (Querying).

5. Feature Extractor

Feature extractors are paired with a data type, for example TermDocumentFrequency is paired with TextData. This means that you can use this feature extractor only for text. This means that extractors are reusable (unlike data readers). Since this is the case let's see how you would use one, and leave the tutorial on writing one for later.

We assume that you have created your problem e.g. scholarships and populated its data with listings from different websites using your Data Reader. Before you start with automatic classification you need to extract features on your data which are later used by the classifier.

```
public FeatureExtractorData performFeatureExtraction ( FeatureExtractor feature_extractor , String
feature_extractor_name , List<Data> data ) {
    FeatureExtractorData feature_extractor_data = new FeatureExtractorData(feature_extractor_name);
    try {
        feature_extractor_data.setFeatures(feature_extractor.getFeatures(data));
    } catch (UnsupportedClassTypeException e) {
        e.printStackTrace();
    }
    feature_extractor_data.setDataset(feature_extractor.getDataset());
    return feature_extractor_data;
}
```

FeatureExtractorData is a data model that wraps all the features we get from a feature extractor on a given list of data. In particular it contains the following information:

`weka.core.Instances dataset` - an instances object containing the descriptions of all the attributes which your extractors measures or calculate. In example, TDF uses terms as attributes and counts their frequency, so the dataset will consist of all possible terms,

`List<Feature> features` - a list of the features, one for each data. A Feature is simply a feature vector (`weka.core.Instance`) and a `data_id(UUID)` to link it to the data from which it was extracted.

Extracting features needs data to operate on which means you will need the `getAllDataForProblem` method which returns a list of all the data. To get a list of data entries that have some property you can run a query on the database. For example to get all data where the level field is present in one would use:

```
DBCcollection coll = dmbgt.getCollection("data");
DBObject query = new BasicDBObject("category_information.level",new
BasicDBObject("$exists","true"));
DBCursor cursor = coll.find(query);
List<Data> data = new ArrayList<Data>(cursor.count());
while ( cursor.hasNext() ) {
    DBObject data_object = cursor.next();
    Data d = databasemgt.getParser().readValue(data_object.toString(), Data.class);
}
```

You can use the same method if you are trying to extract features for a huge list of data, so huge it doesn't fit in memory. You iterate through them with the cursor and extract features on the fly:

```
//sample data
```

```

while ( cursor.hasNext() ) {
    DBObject data_object = cursor.next();
    Data d = databasemgt.getParser().readValue(data_object.toString(), Data.class);
    f_data.addFeature(feature_extractor.extractFeature(d));
}

```

After you have your feature extracotor data you should save it in your DB for later usage

```
databasemgt.addFeatureExtractorDataToProblem(feature_extracotor_data,"scholarships").
```

If you want to add new features (not attributes) to an existing feature extractor data object (probably from some new data you just fetched from the web) you can use the method

```
databasemgt.addFeatureToFeatureExtractorDataForProblem(feature,"TDF","scholarships");
```

6. Writing your own feature extractor

Now we look at how to write your own feature extractor. If you are not interested in this feel free to skip this part. You will first need to choose a feature extraction method. This is a well-studied area for common data domains i.e. textual data or image data. Next create a class and make it extend from FeatureExtractor. You will need to override just one method which is:

```
weka.core.Instance getFeatures(Data d);
```

Note that sometimes to conserve space it is optimal to use weka.core.SparseInstance. Another thing that your feature extractor needs to do is populate the vector attributes with all the possible attributes that can be extracted from a data. The usage of attributes is encapsulated in the FeatureExtractor class and you should not change it. You populate attributes at construction time or while you are extracting features from a given data entry. However, your feature extractor should never add values for attributes not declared previously because then it will invalidate all your hardly-trained classifiers you have carefully stored in database. To enable the second way of populating your attributes you should also override the method:

```
public weka.core.Instance getFeatures(Data d,Instances dataset);
```

in which you will only use the attributes already in the dataset instead of create new ones. It's up to the client to make sure he never invokes the `getFeatures(data)` after the attributes have been put in the database.

```

attributes.addElement(new Attribute("numeric attr"));
attributes.addElement(new Attribute("nominal attr",{ "value1", "value2", "value3" }));

```

Remember what we mentioned above about every feature extractor being paired with a single data type. You should check this in the beginning of your `extractFeatures` method like so.

```

if ( data instanceof TextData ) {
    //all ok, continue with normal execution
} else {
    throw new weka.core.UnssuportedClassTypeException("Expected data type is
:"+TextData.class.toString()+" , passed was data of type :"+data.getClass().toString()+"");
}

```

If you are planning on writing your own feature extractor make sure you check the `weka.core.Instance`, `weka.core.SparseInstance` and `weka.core.Attribute` classes to see how you would use them in your code, or look at how TDF is implemented.

7. Classifiers

Classifiers are general in their use and are not related to any particular data domain. This means that the same classifier can work both on images and on text. You will not usually write classifiers yourself; instead you will probably use the ones that came with Weka. So let's just look at a few basic classifiers operation you would typically use in your application. Before a classifier is used it must be trained using a training dataset. A training dataset is simply a set of data represented by their features i.e. `FeatureExtractorData` for which we have class information, e.g. whether the data belongs in the sports or the politics section as specified on the website from where we originally crawled the data. This is how you would build a decision tree classifier (J48) to classify scholarships listing by their purpose (study or research).

```
Classifier classifier = new J48(); //decision tree
FeatureExtractorData f_data = databasemgt.getFeatureExtractorDataForProblem("TDF","scholarships");
FeatureExtractorDataHelper helper = new FeatureExtractorDataHelper();
Instances dataset = helper.getInstancesAndAddANewClassAttribute("scholarships","purpose", true, null);
classifier.buildClassifier(dataset);
databasemgt.addClassifierToProblem(classifier, classifier_name, problem_name);
```

This code is rather straightforward except for the helper class, so let's elaborate on that. Suppose you've run a TDF feature extractor on your data and you have all your features stored in your database. Now you want to classify on the "purpose" information. To train your classifier you must tell him which of the features belong to which field (class). In Weka this is done by declaring a class attribute. This helper class makes your life easier by saving you the trouble of reading through the Weka documentation and adds the class attribute automatically as long as it is specified in your data when you fetched it from the web. You might wonder why didn't we included this 'class' attribute while we were extracting the features. Well the reason is that the features you get from your feature extractor should be reusable for any field not just "purpose", so as to avoid having to write a different feature extractor for each field.

The last two parameters of the `getInstancesAndAddANewClassAttribute` are boolean `remove_missing_values`, which if set to true will only return the features for any data objects that have the information on that particular field. This is what you would normally use for training the classifier, else if the third parameter is false the method will return the feature vectors for all the data – this is what you will want to use later for the actual classifying.

The second parameter is your category map which if set to null will force the method to find all distinct values of the field (second parameter) and save them in the class attribute.

Once we trained our classifier we want to use it to classify some data:

```
//fetch the classifier and the data from DB
Classifier classifier = databasemgt.getClassifierForProblem("purpose","scholarship");
FeatureExtractorData f_data = databasemgt.getFeatureExtractorDataForProblem("TDF","scholarships");

//add the class attribute using your pre-built category map (this is saved automatically when you
saved your classifier)
FeatureExtractorDataHelper helper = new FeatureExtractorDataHelper();
CategoryMap cm = databasemgt.getCategoryMapForProblem("scholarships");
Instances dataset = helper.getInstancesAndAddANewClassAttribute("scholarships","purpose", false,
cm);
//classify any data object who don't have class information
List<String> classes = new ArrayList<String>();
List<Integer> indices = new ArrayList<Integer>();
for ( int i = dataset.numInstances()-1 ; i >= 0 ; --i) {
    //if an instance is already classified, skip it
    if ( ! dataset.instance(i).hasClassMissing() ) continue;
```

```

        //this is where the actual classification happens
        double class_value = classifier.classifyInstance(dataset.instance(i));
        //look up the class name, from its index
        String class_name = dataset.classAttribute().value((int)class_value);
        classes.add(class_name);
        indices.add(i);
    }
    //update the class information in DB
    for ( int i = 0 ; i < classes.size() ; ++i ) {
        Feature f = f_data.getFeatures().get(indices(i));
        databasemgt.setCategoryInformationForDataForProblem("purpose",classes.get(i),f.getData_id
    ),"scholarships");
    }
}

```

If you want to do this only for a specific set of data you can run a custom query on the database. Another thing you might want to do is to compare the performance of your classifiers. Weka comes with a large pool of classifiers so you need to choose the right one for the job and the only way to do this is test their performance. This can be done using the Evaluation class from Weka

```

Instances train = ...    //from somewhere
Instances test = ...     //from somewhere
// train classifier
Classifier cls = new J48();
cls.buildClassifier(train);
// evaluate classifier and print some statistics
Evaluation eval = new Evaluation(train);
eval.evaluateModel(cls, test);

```

Or if you can't obtain a large amount of data to divide it into training and test partitions you can perform a cross validation.

```

Evaluation eval = new Evaluation(data);
eval.crossValidateModel(cls, data, 10, new Random());

```

This can also be used to adjust the parameters for your classifier.

8. Clustering

Clustering is used to identify identical or near duplicate data entries that would frequently appear on different websites. You can't automatically evaluate your clusterers and you will have to manually inspect the clustering results and see if produced clusters are appropriate. A good choice is to use a hierarchical clusterer where you have more direct control over the clusters creation.

Using a clusterer is simple, but you have to populate the cluster_tasers map as you iterate

```

//create the dataset (list of attributes)
Instances dataset = feature_extractor_data.getInstances();
Map<Integer,ClusterTeaser> cluster_tasers = new HashMap<Integer,ClusterTeaser>();
//build the clusterer
clusterer.buildClusterer(dataset);
//set the cluster for each data object
for ( Feature f : feature_extractor_data.getFeatures() ) {
    //merge your list of attributes (dataset) with the actual values (instance)
    Instance instance = f.getInstance();
    instance.setDataset(dataset);
    //do the actual clustering
    int cluster_number = clusterer.clusterInstance(instance);
    //create a ClusterTeaser for each cluster
    ClusterTeaser cluster_teaser = cluster_tasers.get(cluster_number);
    if ( cluster_teaser == null ) {
        Data data = databasemgt.getDataForProblem(f.getData_id(),"scholarships");
        cluster_teaser = new ClusterTeaser(UUID.randomUUID(), data.getTitle());
        cluster_tasers.put(cluster_number, cluster_teaser);
    }
}

```



```

    }
    databasemgt.setClusterTeaserForDataForProblem(cluster_teaser, f.getData_id(),"scholarships");
}

```

The cluster teaser can be used later when displaying the data on your website.

9. Querying

One of the main purposes of an aggregator is to allow easier access to relevant data. Instead of using a key word search system we will use a more domain specific search using the category map. We will allow the user to choose a value for any of the categories e.g.

Purpose:	Level:	Source:
any	any	any
Study	Bachelor	www.scholarshipportal.com
Research	Master	www.scholarship-links.com
	Phd	www.scholarprhipposition.com

Getting the values that go into the combo boxes is simple

```

//load the category map from database
CategoryMap cm = databasemgt.getCategoryMapForProblem("scholarships");
//find categories that is named purpose and then get all its children
List<<Category>> purpose_sub_categories = cm.getCategory("purpose").getChildren();
//copy over the names of the subcategories
List<String> purpose_values = new ArrayList<String>(purpose_sub_categories.size()+1);
For ( Category c : purpose_sub_categories ) purpose_values.add(c.getname());

```

You can use this procedure for any category even if it is not directly connected to the root category like the Bachelor category in the example above.

After the user has chosen specific values for the categories you can run a query on the database to get the appropriate data, Let us suppose that the user chose purpose: study, source: any level: bachelor -> last-year students.

```

//construct the query
Map<String,String> query = new HashMap<String,String>();
query.put("level", "bachelor");
query.put("bachelor ", "last-year students");
query.put("purpose ", "study");
//fetch the respective data sorted by time of publication
List<Data> result = databasemgt.getDataForQueryForProblemSorted (query,"scholarships","time");

```

You can sort by any value stored in the database just make sure you are familiar with the Jackson naming policy, or inspect the database to see what your field is named. By default the query returns the data that satisfies all given user specifications. It is also easy to run a query using \$or, \$any, \$in if you want more custom queries.

Best of luck with your aggregator!